

gprof: Performance Profiling Guide

1. Introduction and Basic Usage

What is gprof?

Purpose:[81] - gprof is a profiler that monitors performance of your program
- Measures relative performance of functions - Helps detect performance bottlenecks

Why Profile?[81] A function's performance may be poor for two reasons: 1. Each invocation takes too much time 2. The function is called too many times

Three-Step Workflow

Step 1: Compile with -pg flag[81]

```
$ gcc -Wall -pg -o myprog myprog.c
```

Creates executable with profiling instrumentation

Step 2: Run the program[81]

```
$ ./myprog
```

Executes normally and creates gmon.out profile data file

Step 3: Analyze with gprof[81]

```
$ gprof ./myprog
```

Displays flat profile and call graph

2. Understanding Profiling Output

Flat Profile (Timing Profile)

Purpose:[81] Lists functions with detailed profiling information on running times

Example Output:[81]

```
$ gprof -b -p -z ./a.out
```

Flat profile:

Each sample counts as 0.01 seconds.

%	cumulative	self	self	total		
time	seconds	seconds	calls	ns/call	ns/call	name
81.05	0.58	0.58	93324100	6.25	6.25	nextnum
12.69	0.67	0.09	10000000	9.13	61.24	ishappy
4.23	0.71	0.03		-	-	main
0.70	0.71	0.01		-	-	frame_dummy

```
0.00      0.71      0.00      -      -      __do_global_dtors_aux
```

Column Meanings:[81]

Column	Meaning	Example
% time	Percentage of time in this function (excluding called functions)	81.05%
Self seconds	Time spent inside this function only	0.58s
Cumulative seconds	Running total of self times	0.67s
Calls	Number of times function called	93324100
Self ns/call	Average self time per call in nanoseconds	6.25 ns
Total ns/call	Self time plus called functions per invocation	6.25 ns
Name	Function name	nextnum

Call Graph

Purpose:[81] Shows which functions call which functions and call counts

Example Output:[81]

```
index % time  self  children  called  name
[1]   100.0  0.01   0.25    <spontaneous>    main [1]
      1000000/1000000
      ishappy [2]
      1000000/1000000
[2]   96.1   0.03   0.22    1000000    ishappy [2]
      12469250/12469250
      isvisited [3]
      12469250/12469250
      nextnum [4]
```

Reading Call Graph:[81] - [index] is function index - Primary line shows total calls - Above: caller functions - Below: called functions - Format **count1/count2** shows count1 calls out of total count2

3. gprof Options and Commands

Compilation and Execution

Compile with profiling:[81]

```
$ gcc -Wall -pg myprog.c           # Creates a.out
$ gcc -Wall -pg -o myprog myprog.c # Creates myprog
```

Run program with arguments:[81]

```
$ ./a.out arg1 arg2 arg3
```

Display Options

Basic analysis:[81]

```
$ gprof ./a.out           # Full output (flat + call graph)
$ gprof ./a.out gmon.out  # Explicit profile file
```

Option: -b (Compact output)[81]

```
$ gprof -b ./a.out
```

Removes explanatory text, shows only data

Option: -p (Flat profile only)[81]

```
$ gprof -p ./a.out
```

Option: -q (Call graph only)[81]

```
$ gprof -q ./a.out
```

Option: -pfunctionname (Specific function)[81]

```
$ gprof -pnextnum ./a.out
```

Shows flat profile for nextnum function only

Option: -z (Include all functions)[81]

```
$ gprof -z ./a.out
```

Includes functions with zero time and system functions

Combined options:[81]

```
$ gprof -b -p -z ./a.out           # Compact flat profile with all functions
$ gprof -b -pfishappy -z ./a.out   # Compact profile of specific function
```

4. Case Studies

Case Study 1: Happy Numbers

Problem:[81] Check if numbers are happy (repeatedly sum squares of digits until reaching 1 or cycle)

Examples:[81] - 2026 is happy: $2026 \rightarrow 44 \rightarrow 32 \rightarrow 13 \rightarrow 10 \rightarrow 1$ - 2024 is unhappy: $2024 \rightarrow 24 \rightarrow 20 \rightarrow 4 \rightarrow 16 \rightarrow 37 \rightarrow 58 \rightarrow 89 \rightarrow 145 \rightarrow 42 \rightarrow 20$ (cycle)

Four Optimization Attempts:

Attempt 1: Array initialization bottleneck[81]

```
$ gprof -b -p -z ./a.out
%      cumulative      self          self      total
time   seconds        seconds  calls  us/call   us/call   name
99.15   9.32           9.32    100000  93.20    93.20    init
0.11    9.33           0.01   1246773  0.01     0.01    isvisited
```

Problem: init() takes 99.15% initializing large array for each call

Attempt 2: Smaller arrays (math optimization)[81]

```
%      cumulative      self          self      total
time   seconds        seconds  calls  us/call   us/call   name
90.17   1.51           1.51   1000000  1.51     1.51    init
4.84    1.59           0.08   12469340 0.01     0.01    nextnum
```

Improvement: 9.32s \rightarrow 1.69s (5.5 \times faster) Problem: init() still 90.17%

Attempt 3: Dictionary approach[81]

```
%      cumulative      self          self      total
time   seconds        seconds  calls  ns/call   ns/call   name
50.53   0.13           0.13   12469250 10.54    10.54    isvisited
23.32   0.19           0.06   12469250 4.86     4.86     nextnum
11.66   0.22           0.03   1000000  30.32    247.62   ishappy
```

Improvement: 1.69s \rightarrow 0.26s (6.5 \times faster) Problem: isvisited() now bottleneck at 50.53%

Attempt 4: Algorithmic breakthrough[81]

```
%      cumulative      self          self      total
time   seconds        seconds  calls  ns/call   ns/call   name
82.27   0.54           0.54   93324100 5.82     5.82     nextnum
15.38   0.64           0.10   10000000 10.15    58.63    ishappy
```

Key insight: Happy numbers \rightarrow 1, unhappy numbers \rightarrow cycle containing 4 No data structure needed, just check if reaches 1 or 4

Case Study 2: Recursive Fibonacci

Naive Recursion:[81]

```
int Fib(int n) {
    if (n < 0) return -1;
    if (n == 0) return 0;
    if (n == 1) return 1;
    return Fib(n-1) + Fib(n-2);
}
// Call: Fib(32)
```

Call graph output:[81]

index	% time	self	children	called	name
			7049154		
[1]	100.0	0.01	0.00	1+7049154	Fib [1]

Calls: 1 + 7,049,154 recursive calls

With Memoization:[81]

```
int Fib(int n, int F[]) {
    if (F[n] >= 0) return F[n];
    if (n == 0) F[n] = 0;
    else if (n == 1) F[n] = 1;
    else F[n] = Fib(n-1, F) + Fib(n-2, F);
    return F[n];
}
// Call: Fib(32, F)
```

Call graph output:[81]

index	% time	self	children	called	name
			62		
[1]	0.0	0.00	0.00	1+62	Fib [1]

Calls: 1 + 62 recursive calls Reduction: 7,049,154 \rightarrow 62 calls (113,373 \times improvement!)

5. Limitations and Important Notes

Sampling-Based Approach[81]

How sampling works: - gprof samples execution every 0.01 seconds (default)
- Based on samples, makes statistical analysis - Percentages are estimates, not exact

Accuracy Requirements:[81] - Program must run for at least a few seconds for meaningful results - Insufficient samples lead to inaccurate estimates - Sampling

rate cannot be changed

Percentage Limitations:[81] - Percentages may not sum to exactly 100% - Sum may be less than or even larger than 100% - Normal limitation of sampling-based profiling

Functions in Output[81]

Functions not listed: - Functions not called during profiling - Missed all samples - Use `-z` option to include them

Unexpected system functions:[81] - Functions like `frame_dummy`, `__do_global_dtors_aux` - Called by runtime system - Usually account for small percentage

Call Count Notation:[81] - Regular: Single call count - Recursive: `count1+count2` format (count1=non-recursive, count2=recursive) - Caller/called lines: `count1/count2` format

Profiling Limitations[81]

Function-level only: - gprof handles function-level profiling - For line-by-line profiling, use `gcov`

No line-by-line in modern systems:[81] - Line-by-line profiling option `-l` works with old gcc - Recommended to use `gcov` for modern systems