

CS29206 Systems Programming Laboratory

Spring 2024

Introduction to valgrind

Abhijit Das
Pralay Mitra

What is valgrind?

- Valgrind is not value grinder. It is the gate to Valhalla (Hall of the Slain).
- Our valgrind is a memory debugging tool.
- Capable of detecting many common memory-related errors and problems.
- Also used for various memory-related profiling.
- Here, we use only the memcheck feature of valgrind.
 - Memory leaks
 - Invalid memory access

How to use valgrind

- Run as: `valgrind executable <command line options>`
- Example: `valgrind ./a.out 2022 -name "Sad Tijihba"`
- Sample valgrind output

```
==4825== Memcheck, a memory error detector
==4825== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==4825== Using Valgrind-3.15.0 and LibVEX; rerun with -h for copyright info
==4825== Command: ./a.out
==4825==
```

Input/Output of your program

```
==4825==
==4825== HEAP SUMMARY:
==4825==    in use at exit: 0 bytes in 0 blocks
==4825==   total heap usage: 23 allocs, 23 frees, 1,376 bytes allocated
==4825==
==4825== All heap blocks were freed -- no leaks are possible
==4825==
==4825== For lists of detected and suppressed errors, rerun with: -s
==4825== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```

- valgrind can also issue error/warning messages inside your program's transcript.

Common memory issues

Memory leaks

- Memory is allocated in blocks by malloc, calloc, realloc.
- Allocated memory should be free'd when no longer in use.
- Types of memory allocated but not free'd.
 - **Reachable** Blocks not free'd but can be accessed until the end of the program.
 - **Definitely lost** Blocks not free'd and not accessible by any pointer.
 - **Indirectly lost** Blocks not free'd and accessible only by pointers in other lost blocks.
 - **Possibly lost** Blocks not free'd and accessible only by pointers in their interiors.

Invalid memory access

- Read/write a memory location outside the allocated area of an array.

An implementation of sorted linked lists

The node data type

```
typedef struct _node {  
    int data;  
    struct _node *next;  
} node;
```

- In my machine, sizeof(node) is 16.
- We maintain a dummy node at the beginning of the list.
- An initialization function only creates the dummy node.
- An insert function makes a sorted insertion in the list. If the element to be inserted is already present, no new node is created.
- A delete function deletes a data from the list. If that data is not present in the list, the list is kept unchanged.

Delete without freeing: Example

```
node *ldel ( node *L, int x )
{
    node *p;

    p = L;
    while (p -> next) {
        if (p -> next -> data == x) {
            p -> next = p -> next -> next;
            return L;
        }
        if (p -> next -> data > x) break;
        p = p -> next;
    }
    return L;
}
```

Delete without freeing: Example


```
sizeof(node) = 16
insert 5 : 5
insert 14 : 5 14
insert 18 : 5 14 18
insert 4 : 4 5 14 18
insert 13 : 4 5 13 14 18
insert 20 : 4 5 13 14 18 20
insert 8 : 4 5 8 13 14 18 20
insert 9 : 4 5 8 9 13 14 18 20
delete 5 : 4 8 9 13 14 18 20
delete 14 : 4 8 9 13 18 20
delete 18 : 4 8 9 13 20
delete 13 : 4 8 9 20
delete 20 : 4 8 9
==9837==
==9837== HEAP SUMMARY:
==9837==    in use at exit: 144 bytes in 9 blocks
==9837==    total heap usage: 10 allocs, 1 frees, 1,168 bytes allocated
==9837==
==9837== LEAK SUMMARY:
==9837==    definitely lost: 48 bytes in 3 blocks
==9837==    indirectly lost: 32 bytes in 2 blocks
==9837==    possibly lost: 0 bytes in 0 blocks
==9837==    still reachable: 64 bytes in 4 blocks
==9837==    suppressed: 0 bytes in 0 blocks
==9837== Rerun with --leak-check=full to see details of leaked memory
==9837==
==9837== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```

Memory leak example


After all insertions




Delete 5




Delete 14




Memory leak example: Continued




Delete 18



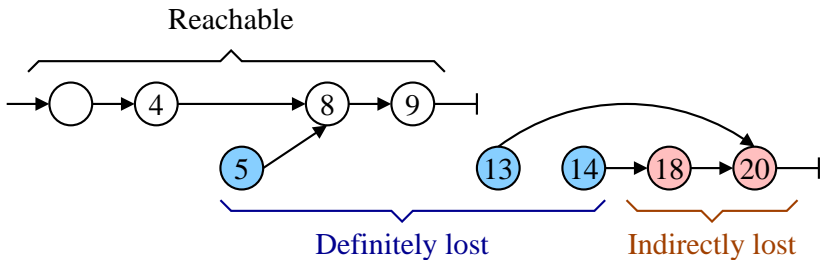
Delete 13



Memory leak example: Continued



Delete 20



Delete with freeing

```
node *ldel ( node *L, int x )
{
    node *p, *q;
    p = L;
    while (p -> next) {
        if (p -> next -> data == x) {
            q = p -> next; p -> next = q -> next; free(q);
            return L;
        }
        if (p -> next -> data > x) break;
        p = p -> next;
    }
    return L;
}
```

```
==10012==
==10012== HEAP SUMMARY:
==10012==      in use at exit: 64 bytes in 4 blocks
==10012==    total heap usage: 10 allocs, 6 frees, 1,168 bytes allocated
==10012==
==10012== LEAK SUMMARY:
==10012==    definitely lost: 0 bytes in 0 blocks
==10012==    indirectly lost: 0 bytes in 0 blocks
==10012==    possibly lost: 0 bytes in 0 blocks
==10012==    still reachable: 64 bytes in 4 blocks
==10012==           suppressed: 0 bytes in 0 blocks
==10012== Rerun with --leak-check=full to see details of leaked memory
==10012==
==10012== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```

Free the reachable nodes at the end

```
void ldestroy ( node *L )
{
    node *p;

    while (L) {
        p = L -> next;
        free(L);
        L = p;
    }
}
```

```
==10160==
==10160== HEAP SUMMARY:
==10160==      in use at exit: 0 bytes in 0 blocks
==10160==    total heap usage: 10 allocs, 10 frees, 1,168 bytes allocated
==10160==
==10160== All heap blocks were freed -- no leaks are possible
==10160==
==10160== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```


Demonstrating possibly lost memory

```
int main ()
{
    int *p, *q;

    p = (int *)malloc(10 * sizeof(int));
    q = p + 5;

    p = (int *)malloc(5 * sizeof(int));

    exit(0);
}
```



Possibly lost memory: valgrind output

```
==4155== Memcheck, a memory error detector
==4155== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==4155== Using Valgrind-3.15.0 and LibVEX; rerun with -h for copyright info
==4155== Command: ./a.out
==4155==
==4155==
==4155== HEAP SUMMARY:
==4155==   in use at exit: 60 bytes in 2 blocks
==4155== total heap usage: 2 allocs, 0 frees, 60 bytes allocated
==4155==
==4155== LEAK SUMMARY:
==4155==   definitely lost: 0 bytes in 0 blocks
==4155==   indirectly lost: 0 bytes in 0 blocks
==4155==   possibly lost: 40 bytes in 1 blocks
==4155==   still reachable: 20 bytes in 1 blocks
==4155==   suppressed: 0 bytes in 0 blocks
==4155== Rerun with --leak-check=full to see details of leaked memory
==4155==
==4155== For lists of detected and suppressed errors, rerun with: -s
==4155== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```

Overflow in arrays

```
#include <stdio.h>
#include <stdlib.h>

int main ()
{
    int n = 16, i, *A;

    A = (int *)malloc(n * sizeof(int));
    printf("A starts at %p, and ends at %p\n", A, A+n-1);
    for (i=1; i<=n; ++i) A[i] = i * i;
    for (i=1; i<=n; ++i) printf("%d ", A[i]);
    printf("\n");

    free(A);

    exit(0);
}
```

Overflow in arrays: valgrind output

```
==13180== Memcheck, a memory error detector
==13180== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==13180== Using Valgrind-3.15.0 and LibVEX; rerun with -h for copyright info
==13180== Command: ./a.out
==13180==
A starts at 0x4a5a040, and ends at 0x4a5a07c
==13180== Invalid write of size 4
==13180==    at 0x109240: main (in /home/abhij/IITKGP/course/lab/SPL/Spring22/prog/valgrind/a.out)
==13180== Address 0x4a5a080 is 0 bytes after a block of size 64 alloc'd
==13180==    at 0x483B7F3: malloc (in /usr/lib/x86_64-linux-gnu/valgrind/vgpreload_memcheck-amd64-linux.so)
==13180==    by 0x1091EC: main (in /home/abhij/IITKGP/course/lab/SPL/Spring22/prog/valgrind/a.out)
==13180==
==13180== Invalid read of size 4
==13180==    at 0x10926B: main (in /home/abhij/IITKGP/course/lab/SPL/Spring22/prog/valgrind/a.out)
==13180== Address 0x4a5a080 is 0 bytes after a block of size 64 alloc'd
==13180==    at 0x483B7F3: malloc (in /usr/lib/x86_64-linux-gnu/valgrind/vgpreload_memcheck-amd64-linux.so)
==13180==    by 0x1091EC: main (in /home/abhij/IITKGP/course/lab/SPL/Spring22/prog/valgrind/a.out)
==13180==
1 4 9 16 25 36 49 64 81 100 121 144 169 196 225 256
==13180==
==13180== HEAP SUMMARY:
==13180==    in use at exit: 0 bytes in 0 blocks
==13180==   total heap usage: 2 allocs, 2 frees, 1,088 bytes allocated
==13180==
==13180== All heap blocks were freed -- no leaks are possible
==13180==
==13180== For lists of detected and suppressed errors, rerun with: -s
==13180== ERROR SUMMARY: 2 errors from 2 contexts (suppressed: 0 from 0)
```


Buffer overflow: Example

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main ( int argc, char *argv[] )
{
    char *wnote = malloc(32);

    if (argc == 1) exit(1);

    printf("The input has size %ld\n", strlen(argv[1]));
    printf("wnote starts at %p and ends at %p\n", wnote, wnote + 31);
    sprintf(wnote, "Welcome to %s", argv[1]);
    printf("%s\n", wnote);

    free(wnote);

    exit(0);
}
```

- Compile and run as follows. There may be no errors.

```
./a.out "Systems Programming Laboratory"
```

- Run as follows to see the problems.

```
valgrind ./a.out "Systems Programming Laboratory"
```

Buffer overflow: valgrind output

```
==12432== Command: ./a.out Systems\ Programming\ Laboratory
==12432==
The input has size 30
wnote starts at 0x4a5a040 and ends at 0x4a5a05f
==12432== Invalid write of size 1
==12432== Address 0x4a5a060 is 0 bytes after a block of size 32 alloc'd
==12432== Invalid write of size 1
==12432== Address 0x4a5a069 is 9 bytes after a block of size 32 alloc'd
==12432== Invalid read of size 1
==12432== Address 0x4a5a060 is 0 bytes after a block of size 32 alloc'd
==12432== Invalid read of size 1
==12432== Address 0x4a5a068 is 8 bytes after a block of size 32 alloc'd
==12432== Invalid read of size 1
==12432== Address 0x4a5a066 is 6 bytes after a block of size 32 alloc'd
==12432==
Welcome to Systems Programming Laboratory
==12432==
==12432== HEAP SUMMARY:
==12432==      in use at exit: 0 bytes in 0 blocks
==12432==    total heap usage: 2 allocs, 2 frees, 1,056 bytes allocated
==12432==
==12432== All heap blocks were freed -- no leaks are possible
==12432==
==12432== For lists of detected and suppressed errors, rerun with: -s
==12432== ERROR SUMMARY: 38 errors from 6 contexts (suppressed: 0 from 0)
```